

Deep Learning

2. Convolutional Neural Networks

Term 4, 2025

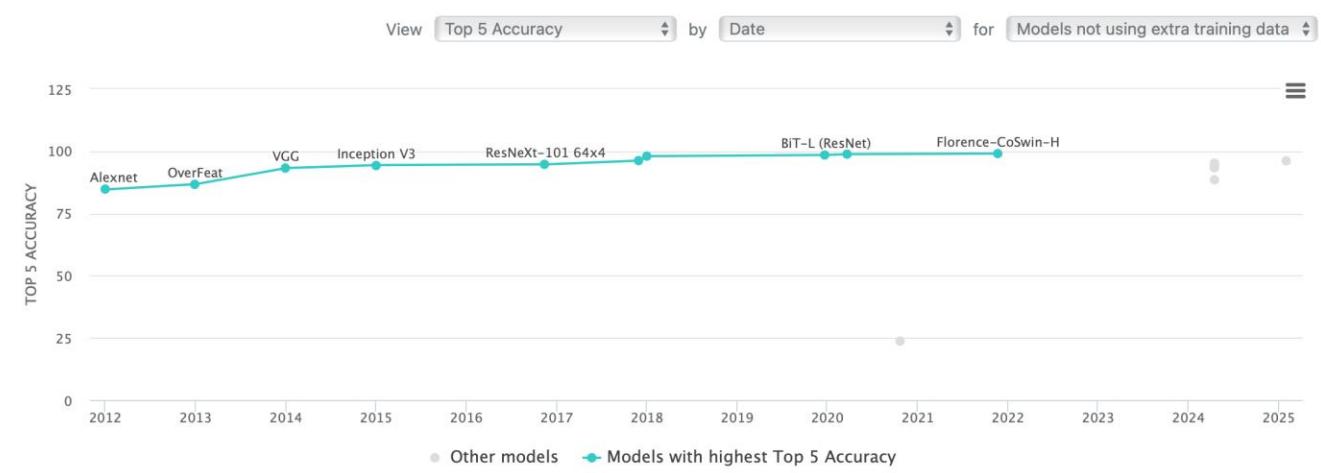
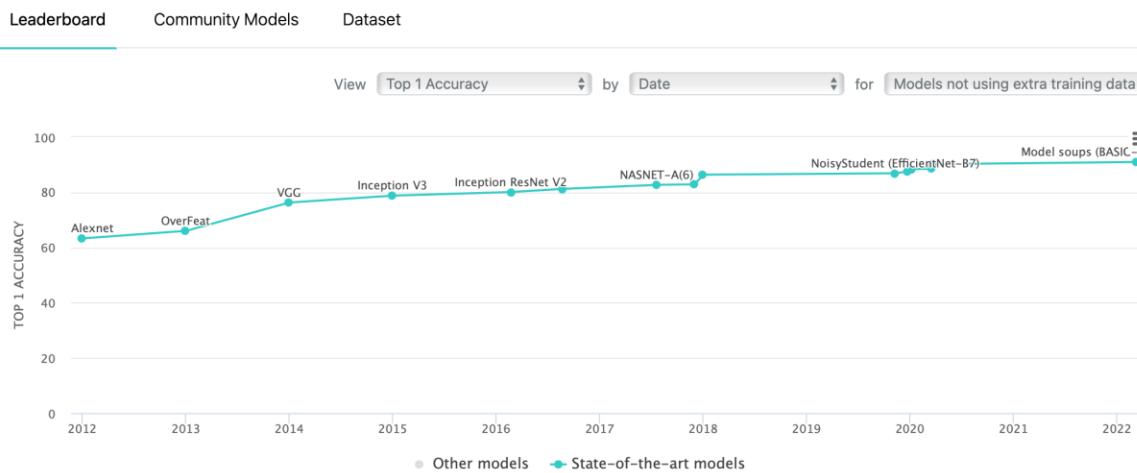
Skoltech

Outline of the lecture

- Introduction
- Biological inspiration
- Foundations in classical computer vision
- Why convolutions work for images
- Mathematical definition
- Early beginnings of CNNs
- The CNN winter and revival
- The deep learning revolution

Introduction

- CNNs revolutionized computer vision:
 - They became the dominant approach for image recognition tasks since 2012
 - They achieved superhuman performance on many visual tasks
 - Very high ImageNet classification accuracy



<https://paperswithcode.com/sota/image-classification-on-imagenet>

Introduction

- Inspired by the visual cortex of animals
 - The hierarchical structure of CNNs mirrors the organization of the mammalian visual system
 - Early layers detect simple features like edges and textures, while deeper layers recognize complex patterns
 - Hubel and Wiesel's Nobel Prize-winning work on the visual cortex provided the biological foundation
 - <https://PMC1359523>/
 - <https://www.nobelprize.org/prizes/medicine/1981/press-release/>

Introduction

- CNNs are designed to process grid-like data (images)
 - Unlike fully connected networks, CNNs preserve spatial relationships in data
 - CNNs can efficiently process high-dimensional visual data with fewer parameters
 - Mathematical properties of convolutions make them particularly suitable for image processing
 - <https://arxiv.org/abs/1803.01164>

Introduction

- Bridge between classical signal processing and modern deep learning
 - CNNs incorporate principles from traditional signal processing (filtering, downsampling)
 - CNNs automate the feature engineering process that was previously done manually
 - Zeiler and Fergus visualized learned features (we will see them later)
 - <https://arxiv.org/abs/1311.2901>

Biological Inspiration

- Hubel and Wiesel's experiments (1959-1962)
 - David Hubel and Torsten Wiesel recorded neural activity in the cat's visual cortex
 - They discovered that neurons respond to specific patterns of light in their receptive fields
 - Their 1962 paper established the foundation for understanding visual processing
 - <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC1359523/>
- They received the Nobel Prize in Physiology or Medicine in 1981 for this work
 - <https://www.nobelprize.org/prizes/medicine/1981/press-release/>

Biological Inspiration

- Discovery of simple and complex cells in visual cortex
 - Simple cells respond to oriented edges at specific positions
 - Complex cells respond to oriented edges regardless of position
 - This hierarchical organization inspired the design of CNNs
 - Fukushima's Neocognitron (1980) was the first model to implement this architecture
 - <https://www.cs.princeton.edu/courses/archive/spr08/cos598B/Readings/Fukushima1980.pdf>

Biological Inspiration

- Cells respond to specific patterns within receptive fields
 - Each neuron has a limited "view" of the visual field (receptive field)
 - Neurons at different levels have different receptive field sizes
 - This property is implemented in CNNs through varying kernel sizes and network depth
 - Visualization of receptive fields in CNNs shows striking similarities to biological systems
 - <https://distill.pub/2017/feature-visualization>
 - Will be displayed later

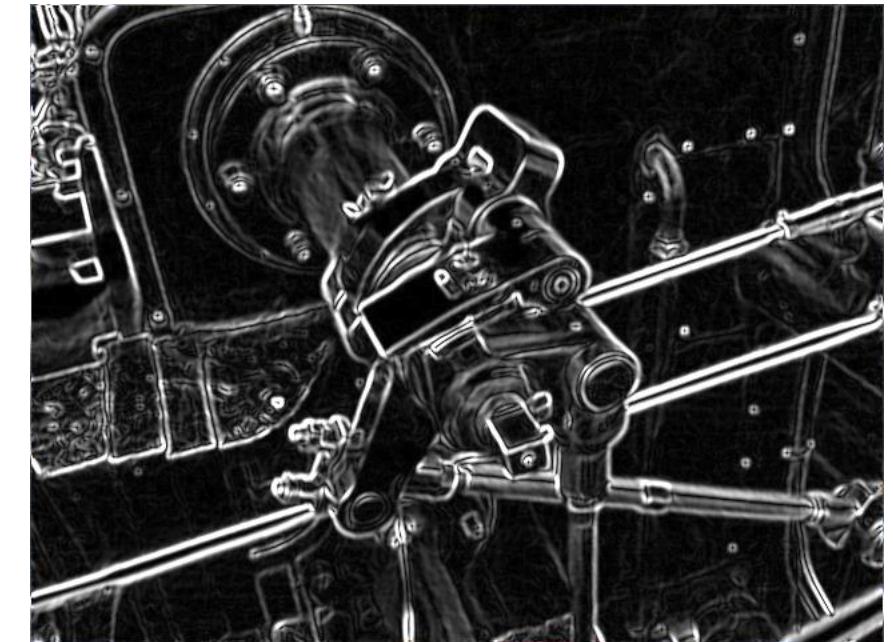
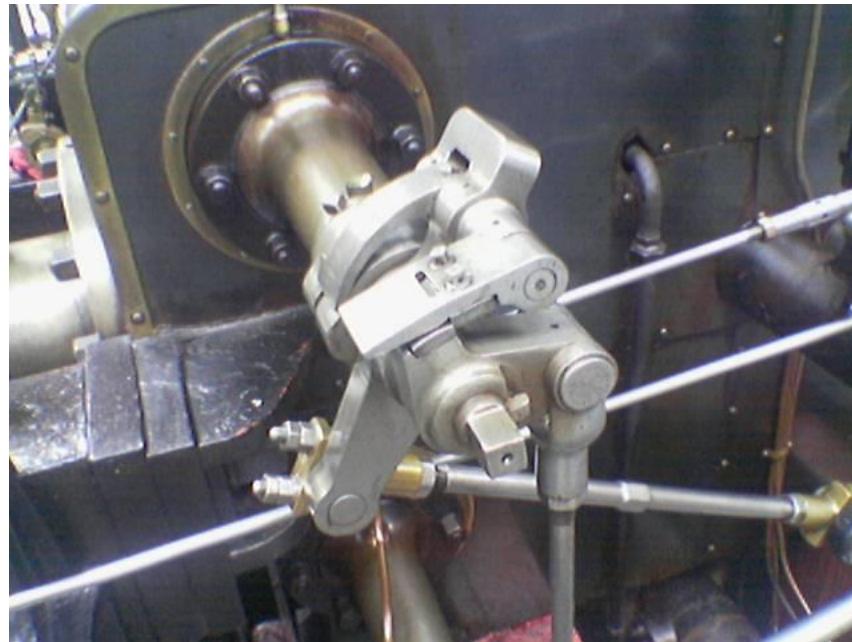
Biological Inspiration

- Hierarchical processing of visual information
 - Information flows from simple features to complex patterns
 - Each layer builds upon the representations from previous layers
 - This hierarchical structure enables CNNs to learn increasingly abstract features
 - Zeiler and Fergus's visualization technique (2013) revealed what each layer learns
 - <https://arxiv.org/abs/1311.2901>
 - Will be displayed later

Foundations in Classical Computer Vision

- Traditional CV used hand-crafted filters
 - Edge detection (Sobel)

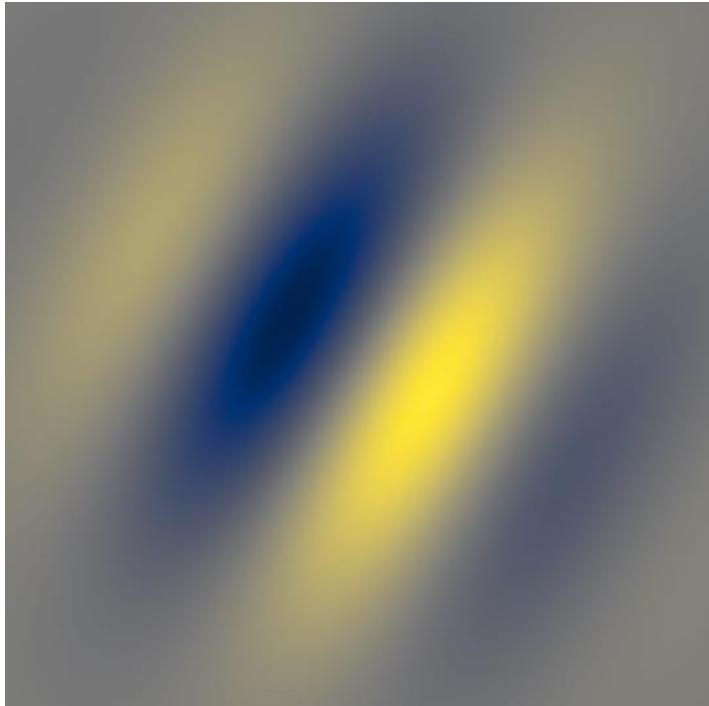
$$\mathbf{G}_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix} * \mathbf{A}$$
$$\mathbf{G}_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix} * \mathbf{A}$$



By Simpsons contributor, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=8904364>

Foundations in Classical Computer Vision

- Traditional CV used hand-crafted filters
 - Gabor filters



By AkanoToE - Own work based on: Gabor filter.png, CC BY-SA 4.0, <https://commons.wikimedia.org/w/index.php?curid=88998601>

By MrJacobs – Own work, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=23580768>

Foundations in Classical Computer Vision

- Traditional CV used hand-crafted filters
 - Sharpening
 - Sharpening enhances edges by subtracting a blurred version from the original
 - Example kernel: $[[0, -1, 0], [-1, 5, -1], [0, -1, 0]]$
 - https://en.wikipedia.org/wiki/Unsharp_masking



By Nevit Dilmen - Own work, Public Domain, <https://commons.wikimedia.org/w/index.php?curid=3884055>

Foundations in Classical Computer Vision

- Traditional CV used hand-crafted filters
 - Strong edge detection
 - Laplacian of Gaussian (LoG) detects edges at multiple scales
 - Gaussian smoothing + Laplacian (sum of 2nd order derivatives) filters
 - <https://homepages.inf.ed.ac.uk/rbf/HIPR2/log.htm>

Why Convolutions Work for Images

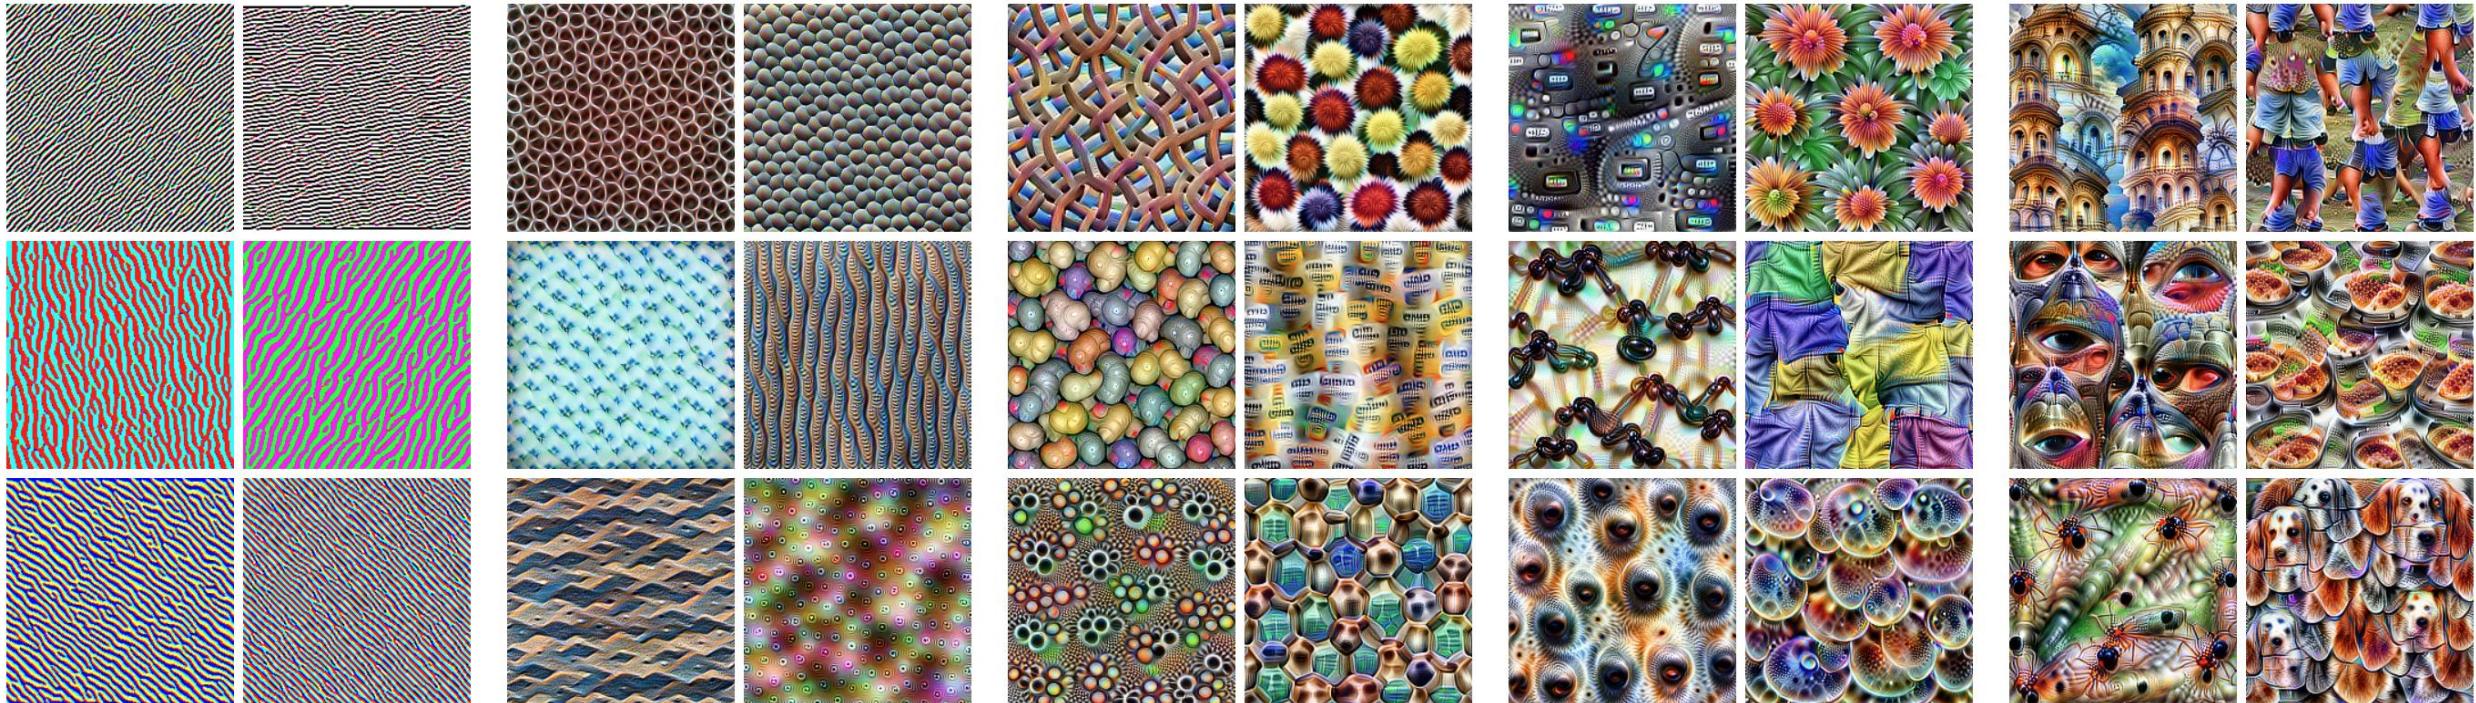
- Local connectivity
 - CNNs process small patches of the image at each layer
 - This reduces the number of parameters compared to fully connected networks
 - For a 1000×1000 image with a single input channel, a fully connected layer would need 10^{12} parameters, while a CNN with 3×3 kernels needs only 9 parameters per filter
 - The same weights are used for all positions in the image
 - This enables translation invariance and further reduces parameters
 - A 3×3 filter with 64 channels has only $3 \times 3 \times 64 = 576$ parameters, regardless of image size

Why Convolutions Work for Images

- Spatial hierarchy
 - Deeper layers capture more complex patterns
 - Early layers detect simple features (edges, textures)
 - Middle layers detect parts (eyes, wheels)
 - Deep layers detect objects (faces, cars)
 - Visualization of hierarchical features
 - <https://distill.pub/2017/feature-visualization/>

Why Convolutions Work for Images

- Spatial hierarchy: Deeper layers capture more complex patterns
 - Feature visualization allows us to see how GoogLeNet



Why Convolutions Work for Images

- Translation invariance
 - Detect features regardless of position
 - CNNs can recognize objects even if they appear in different locations
 - This is achieved through the combination of convolutions and pooling
- Images exhibit locality properties
 - Nearby pixels are more related than distant ones
 - This property makes convolutions a natural choice for image processing
 - It's why CNNs are more efficient than fully connected networks for images

Why Convolutions Work for Images

- Cross-correlation vs. Convolution
 - In practice, CNNs use cross-correlation, not true convolution
 - In true convolution, the kernel is flipped before applying
 - CNNs typically use cross-correlation (no flipping)
 - This is a minor difference that doesn't affect the learning process
 - Convolution is actually used in the backward propagation phase
 - <https://en.wikipedia.org/wiki/Cross-correlation>

Mathematical Definition

- Mathematical definition of a 2D convolution

$$(I * K)(x, y) = \sum_{i=-a}^a \sum_{j=-b}^b I(x - i, y - j)K(i, j)$$

- The learnable kernel K defines the weights of convolution
- I is the input image or feature map
- (x, y) are the coordinates in the output feature map

Mathematical Definition

- Mathematical definition of a 2D cross correlation

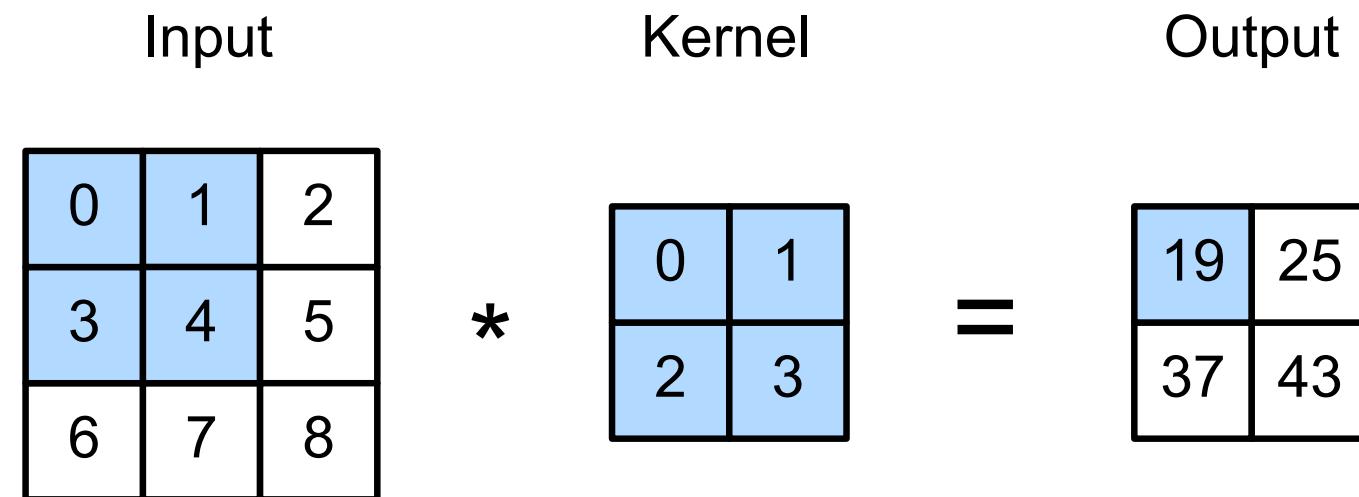
$$(I * K)(x, y) = \sum_{i=-a}^a \sum_{j=-b}^b I(x + i, y + j)K(i, j)$$

- The learnable kernel K defines the weights of cross correlation
- I is the input image or feature map
- (x, y) are the coordinates in the output feature map

Mathematical Definition

- Mathematical definition of a 2D cross correlation

$$(I * K)(x, y) = \sum_{i=-a}^a \sum_{j=-b}^b I(x + i, y + j)K(i, j)$$



Mathematical definition

- Cross correlation

$$O(x, y) = \sum_{i=-a}^a \sum_{j=-b}^b I(s_x x + d_x i, s_y y + d_y j) K(i, j)$$

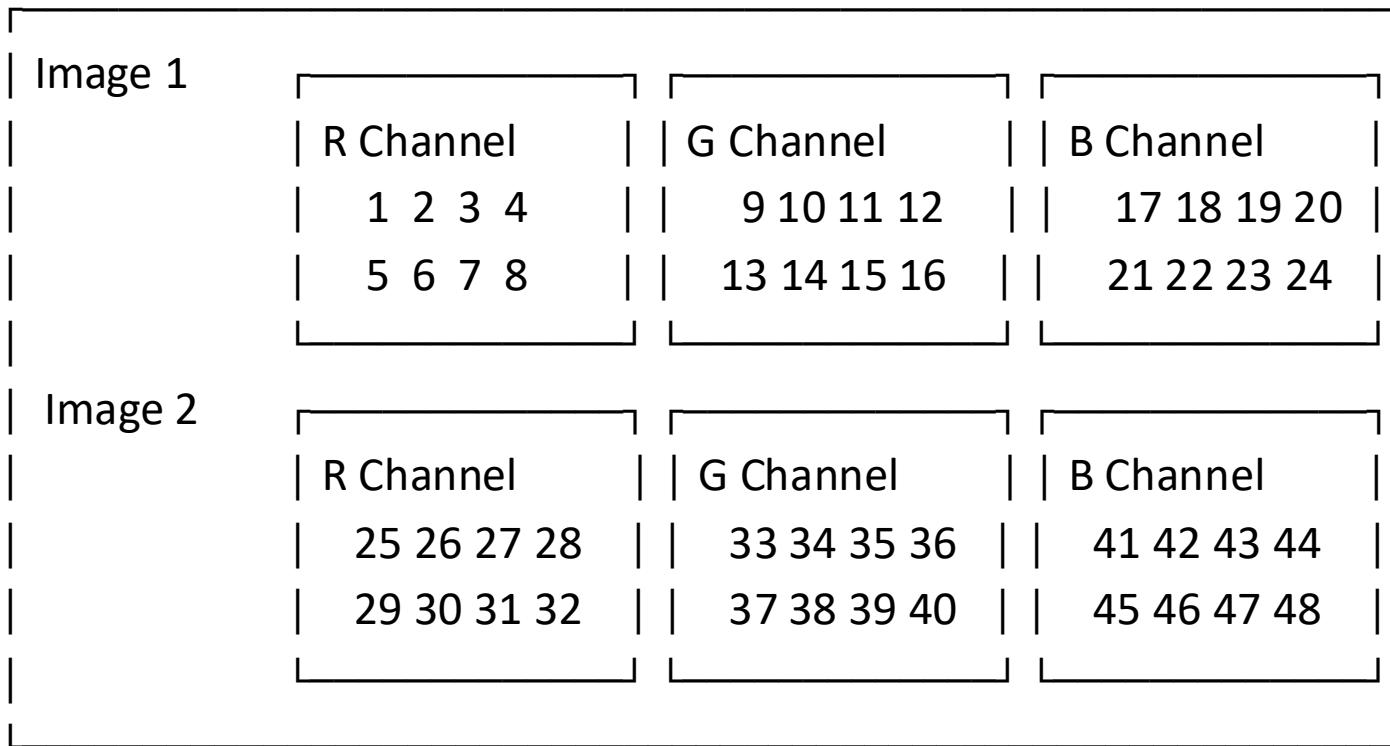
- Padding expands bounds of an input
 - Extension is filled with zeros or by some other rule
- Strides (s_x, s_y)
 - Defines size of a step in each direction of output
- Dilation (d_x, d_y)
 - Defines size of a step in each direction of kernel
- Visualised at
https://github.com/vdumoulin/conv_arithmetic/blob/master/README.md

Data Formats

- Shape of a batch of images is described by
 - N: Batch size (number of images)
 - C: Number of channels (e.g., 3 for RGB)
 - H: Height of the image
 - W: Width of the image
- Batches are stored in one of the standard row-major formats
 - NHWC: better for CPUs and Tensor-core GPUs, default for Tensorflow
 - NCHW: better for non-Tensor-core GPUs, default for Pytorch
 - <https://forums.developer.nvidia.com/t/nhwc-vs-nchw-convolution/111065>

NCHW Data Format

- NCHW Format (2 images, 3 channels, 2x4 pixels):



Forward pass of a Convolutional Layer

- Input
 - Feature map of shape $H \times W \times C_i$
 - H and W are the height and width of the input
 - C_i is the number of input channels (e.g., 3 for RGB images)
- Kernel
 - Shape is $H_K \times W_K \times C_i \times C_o$
 - H_K and W_K are the height and width of the kernel
 - C_i is the number of input channels
 - C_o is the number of output channels (number of filters)

Forward pass of a Convolutional Layer

- Output

$$O(n, m, c_o) = \sum_{i, j, c_o} [I(n + i, m + j, c_i) \times K(i, j, c_i, c_o)] + b(c_o)$$

- Feature map of shape $H' \times W' \times C_o$
- H' and W' are the height and width of the output
 - Depend on input size, kernel size, padding, stride and dilation
 - What is the exact dependency?
- C_o is the number of output channels (number of filters)
- $b(c_o)$ is the bias term for each output channel

Backward pass of a Convolutional Layer

- Gradient w.r.t. input:

$$\frac{\partial L}{\partial I(n, m, c_i)} = \sum_{i,j,c_o} \left[\frac{\partial L}{\partial O(n - i, m - j, c_o)} \times K(i, j, c_i, c_o) \right]$$

- This is a convolution
- This formula computes how the loss changes with respect to each input element
- It is used for backpropagation through the network

Backward pass of a Convolutional Layer

- Gradient w.r.t. kernel:

$$\frac{\partial L}{\partial K(i, j, c_i, c_o)} = \sum_{n,m} \left[\frac{\partial L}{\partial O(n, m, c_o)} \times I(n + i, m + j, c_i) \right]$$

- This formula computes how the loss changes with respect to each kernel element
- It is used to update the kernel weights during training

Backward pass of a Convolutional Layer

- Gradient w.r.t. bias:

$$\frac{\partial L}{\partial b(c_o)} = \sum_{n,m} \frac{\partial L}{\partial O(n, m, c_o)}$$

- This formula computes how the loss changes with respect to each bias element
- It is used to update the convolutional bias during training

Early beginnings of CNNs

- Kunihiko Fukushima's Neocognitron (1980)
 - Hierarchical, multi-layered network
 - First model with a hierarchical structure similar to the visual cortex
 - <https://www.cs.princeton.edu/courses/archive/spr08/cos598B/Readings/Fukushima1980.pdf>
 - First implementation of convolutional structure
 - Used local receptive fields and shared weights
 - Pre-dated modern CNNs by almost two decades
 - Self-organizing model inspired by visual cortex
 - Could learn to recognize patterns without explicit supervision
 - Based on biological principles of visual processing

Early beginnings of CNNs

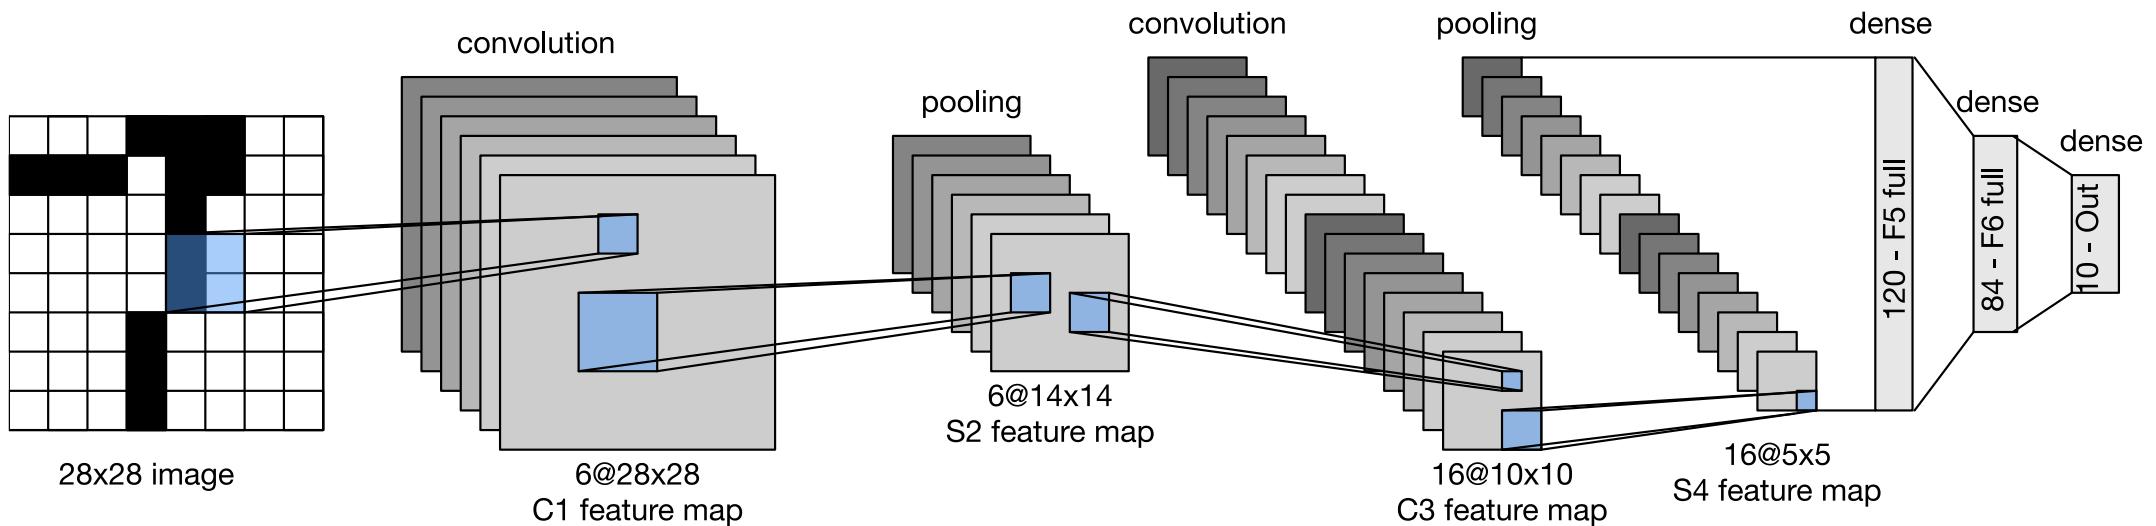
- Yann LeCun's LeNet-5 (1998)
 - First modern CNN architecture
 - http://vision.stanford.edu/cs598_spring07/papers/Lecun98.pdf
 - Handwritten digit recognition (MNIST)
 - Achieved 0.8% error rate on the MNIST dataset
 - Used by banks to recognize handwritten digits on checks
 - Used backpropagation for training
 - Demonstrated that CNNs could be trained end-to-end
 - Showed the power of gradient-based learning
 - Used average pooling for downsampling

Early beginnings of CNNs

- What is an average pooling?
 - Pooling is a common operation in CNNs used to downsample feature maps
 - Average pooling takes the average value of each sub-region of the input
 - It is a special case of a convolutional layer with predefined kernel
 - It reduces the size of the feature map while introducing a form of spatial invariance
 - Average pooling helps to reduce overfitting and improve generalization
 - However, it may lose information about the precise location of features
 - Alternatives to average pooling include **max pooling**, **adaptive pooling**
 - Size of the **adaptive pooling** is defined by the input and output shapes

Early beginnings of CNNs

- Yann LeCun's LeNet-5 (1998)



Early beginnings of CNNs

- Yann LeCun's LeNet-5 (1998)
 - Trained and tested on MNIST dataset (60,000 samples)
 - Short for "Modified National Institute of Standards and Technology database "
 - MNIST became the "Hello World" of deep learning
 - "Everything works on MNIST"
 - <https://yann.lecun.org/exdb/mnist/index.html>

The CNN Winter and Revival

- CNNs faced limited adoption (1990s-2000s)
 - Computational limitations
 - Training deep networks required significant computational resources
 - GPUs were not yet widely used for deep learning
 - Limited training data
 - Small datasets like MNIST were insufficient for deep networks
 - Data augmentation techniques were not yet developed
 - SVM and other methods dominated CV
 - Support Vector Machines and other shallow models were more popular
 - They performed well on small datasets and were easier to train

The CNN Winter and Revival

- ImageNet
 - 2006: Fei-Fei Li proposed the creation of ImageNet
 - "While most people pay attention to models, let's pay attention to data"
 - July 2008: ImageNet had zero images
 - December 2008: 3 million images across 6000 synsets
 - April 2010: 11 million images across 15000 synsets
 - Today (March 2025): 14.2 million images across 21841 synsets
 - Made possible through crowdsourcing on Amazon's Mechanical Turk
 - Provided enough data to train deep networks
 - <https://image-net.org/>

The CNN Winter and Revival

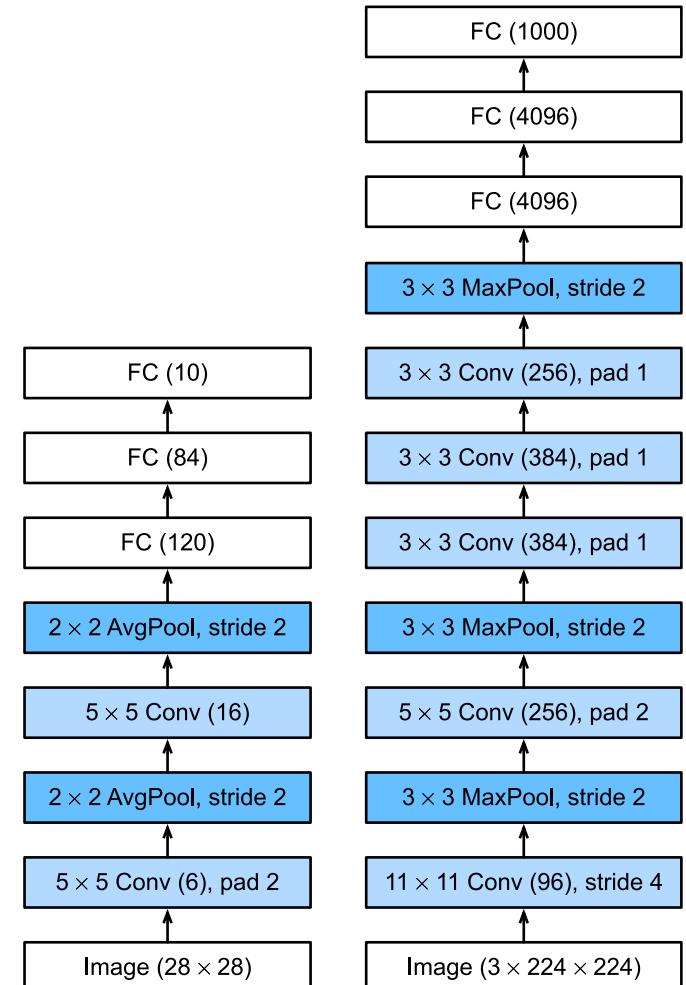
- ILSVRC (ImageNet Large Scale Visual Recognition Competition)
 - Organized in 2010 for the first time
 - Provided an ImageNet large scale subset for training
 - 1.2 million training images across 1,000 categories
 - Annual competition for image classification (2010-2017)
 - ILSVRC became the benchmark for computer vision
 - <https://image-net.org/challenges/LSVRC/>

The AlexNet Breakthrough

- The first deep CNN to win ILSVRC
 - Achieved 15.3% top-5 error rate (previous best: 26%)
 - This breakthrough sparked the deep learning revolution
 - <https://image-net.org/challenges/LSVRC/2012/results.html>
 - https://papers.nips.cc/paper_files/paper/2012/file/c399862d3b9d6b76c8436e924a68c45b-Paper.pdf
- Key innovations
 - Simple ReLU activation: $f(x) = \max(x, 0)$
 - Dropout: zeros random values during training
 - GPU training: 5-6 days to train on two GTX 580 GPUs with 3GB VRAM

The AlexNet Breakthrough

- Architecture limitations
 - Large convolutional filters in early layers
 - 11×11 and 5×5 filters in early layers
 - Inefficient compared to smaller filters
 - Two large MLP layers at the end
 - 6400×4096 and 4096×4096 linear layers
 - These fully connected layers contain most of the parameters
 - Inefficient for modern CNN architectures
 - Total of around 60M parameters
 - Large number of parameters makes the model computationally expensive



The AlexNet Breakthrough

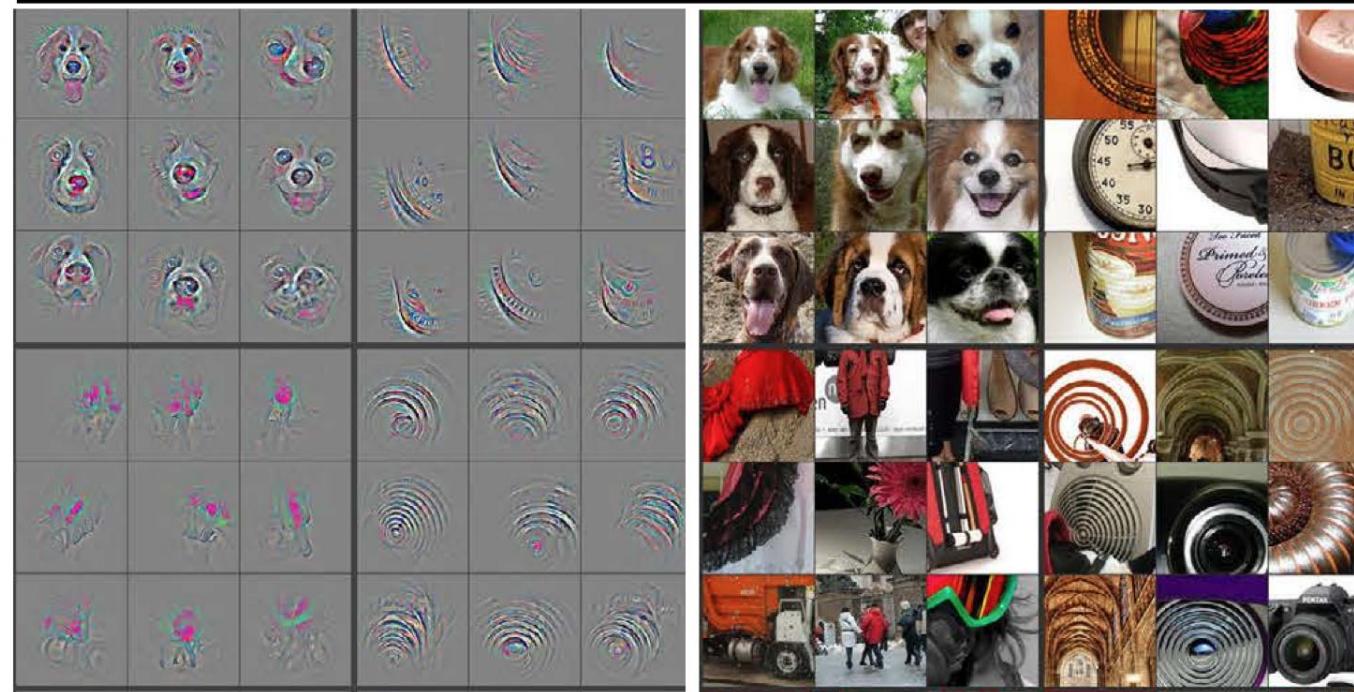
- Visually presented 96 convolutional kernels of size $11 \times 11 \times 3$ learned by the first convolutional layer on the $224 \times 224 \times 3$ input images



https://papers.nips.cc/paper_files/paper/2012/file/c399862d3b9d6b76c8436e924a68c45b-Paper.pdf

The Deep Learning Revolution

- ZFNet (2013): Visualization of CNN layers
 - Improved AlexNet by adjusting hyperparameters
 - Developed visualization techniques to understand what CNNs learn
 - <https://arxiv.org/abs/1311.2901>

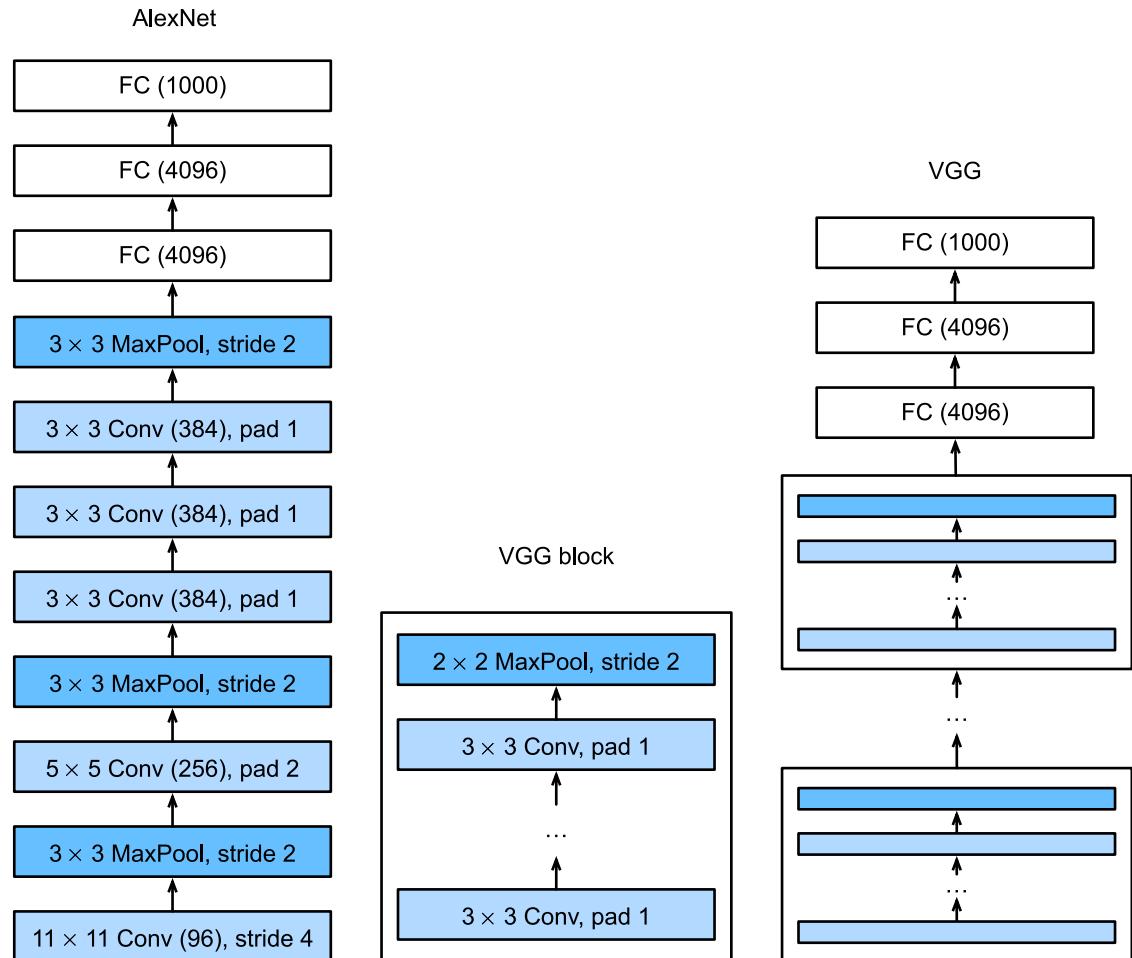


The Deep Learning Revolution

- VGGNet (2014): Simplicity and depth (3x3 convolutions)
 - Won ILSVRC 2014
 - Key idea: Use multiple 3×3 convolutions between MaxPooling downsampling
 - Example: two successive 3x3 convolutional layers against a single 5x5 layer
 - Receptive field is the same and output feature maps are of the same shape
 - Two 3x3 layers require 18 parameters
 - Single 5x5 layer requires 25 parameters
 - Showed that deep and narrow convolutions outperform wider counterparts
 - Replaces larger filters with stacks of 3×3 filters
 - 3 x 3 convolutions became de-facto standard.
 - <https://arxiv.org/abs/1409.1556>

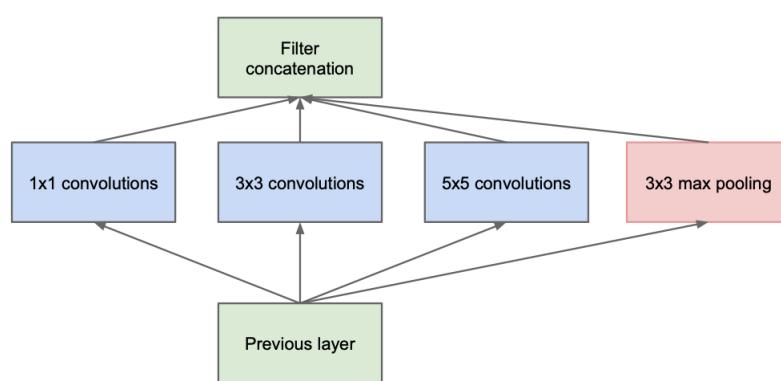
The Deep Learning Revolution

- VGGNet (2014)
 - VGG networks trained on ImageNet are excellent **feature extractors**.
 - VGG-11: 8 convolutional and 3 fully-connected layers
 - VGG-16: 13 convolutional and 3 fully-connected layers
 - VGG-19: 16 convolutional and 3 fully-connected layers

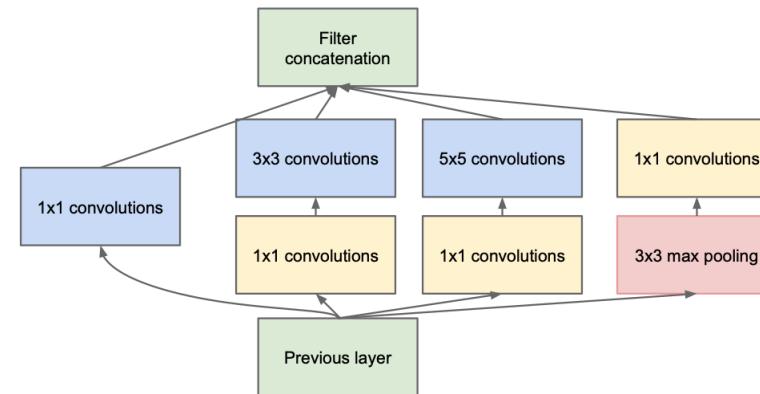


The Deep Learning Revolution

- GoogLeNet/Inception (2014): Inception modules
 - Won ILSVRC 2014
 - Multi-branch networks
 - Parallel paths with different filter sizes
 - <https://arxiv.org/abs/1409.4842>



(a) Inception module, naïve version



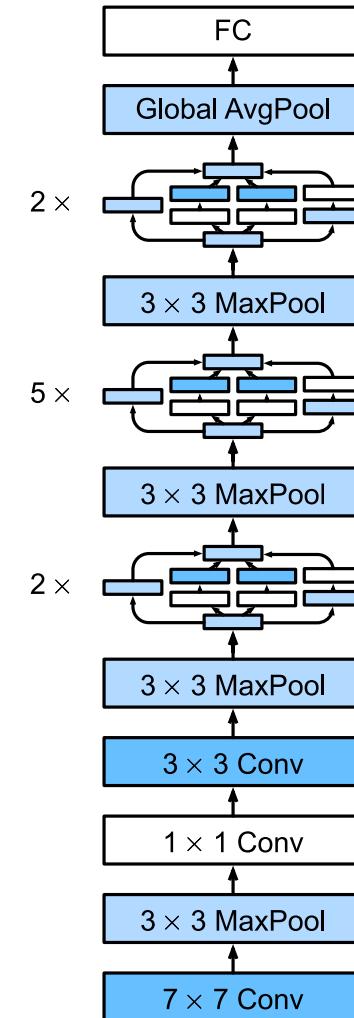
(b) Inception module with dimension reductions

The Deep Learning Revolution

- GoogLeNet/Inception (2014)
 - First few layers: feature extracting stem
 - Last layer: classification head
 - Everything in between: body

Team	Year	Place	Error (top-5)	Uses external data
SuperVision	2012	1st	16.4%	no
SuperVision	2012	1st	15.3%	Imagenet 22k
Clarifai	2013	1st	11.7%	no
Clarifai	2013	1st	11.2%	Imagenet 22k
MSRA	2014	3rd	7.35%	no
VGG	2014	2nd	7.32%	no
GoogLeNet	2014	1st	6.67%	no

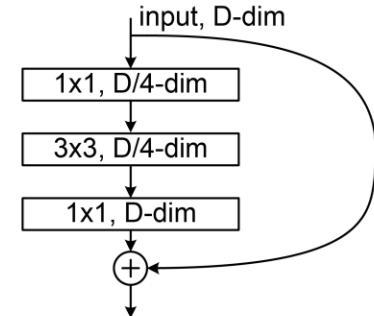
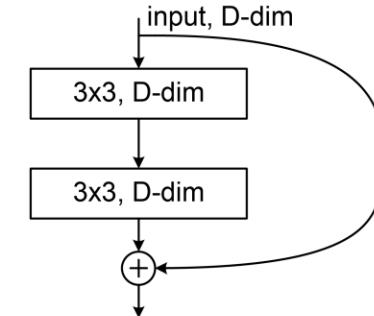
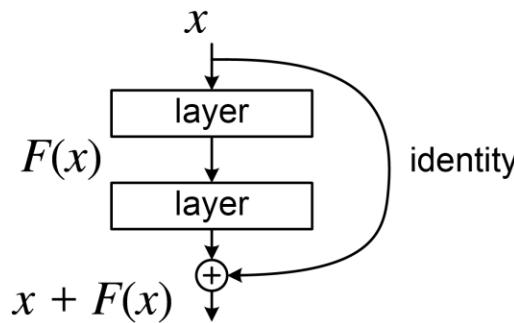
<https://arxiv.org/pdf/1409.4842>



By Zhang, Aston and Lipton, Zachary C. and Li, Mu and Smola, Alexander J. -
<https://github.com/d2l-ai/d2l-en>, CC BY-SA 4.0,
<https://commons.wikimedia.org/w/index.php?curid=152265655>

The Deep Learning Revolution

- ResNet (2015): Skip connections, ultra-deep networks
 - Residual block: Instead of mapping $y := f(x)$, learn $y := x + f(x)$
 - Backprop without residual: $\frac{\partial L}{\partial x} = \frac{\partial L}{\partial y} \cdot \frac{\partial y}{\partial x}$
 - Backprop with residual: $\frac{\partial L}{\partial x} = \frac{\partial L}{\partial y} \cdot \frac{\partial y}{\partial x} + \frac{\partial L}{\partial y}$

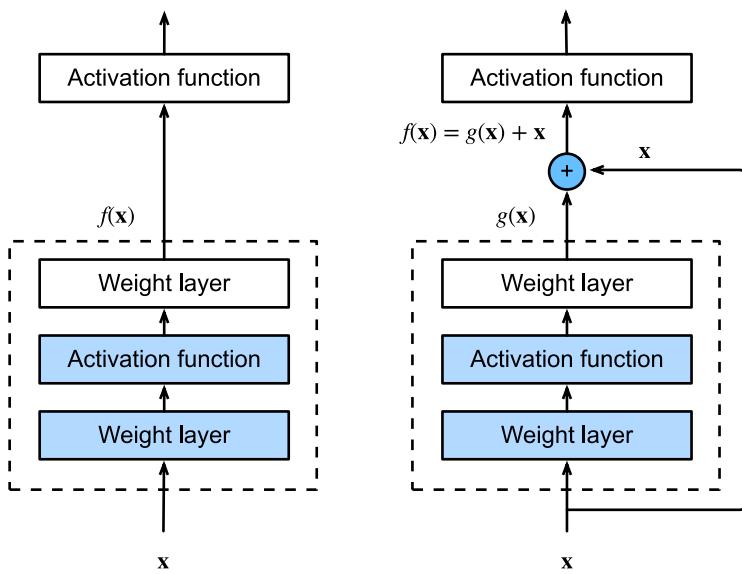


The Deep Learning Revolution

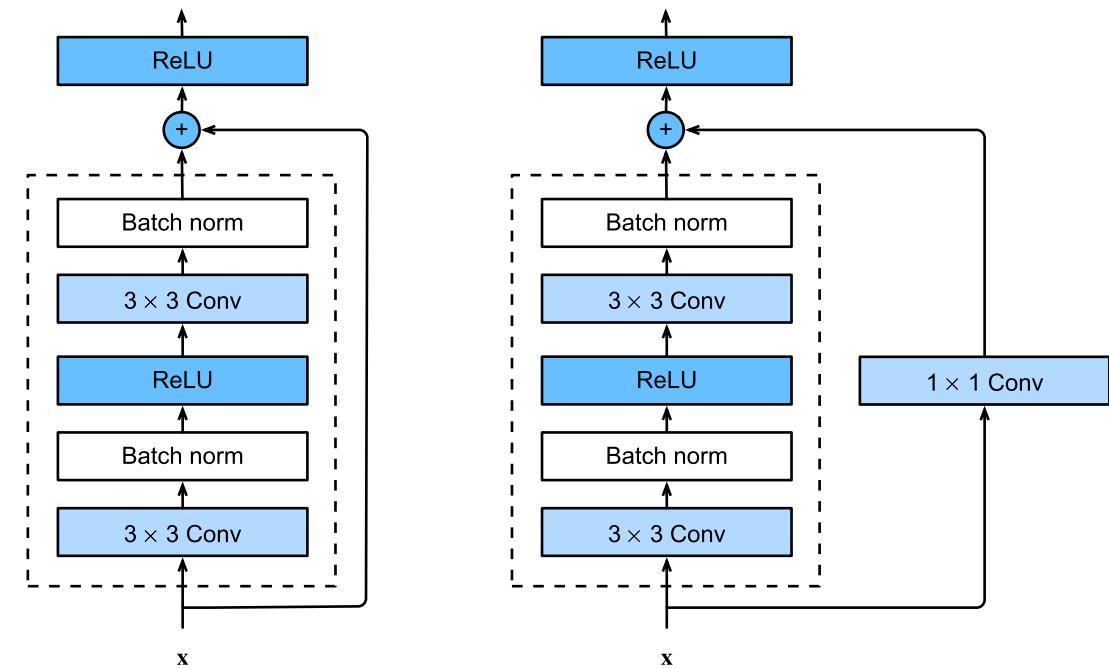
- ResNet (2015): Skip connections, ultra-deep networks
 - Residual block: Instead of mapping $y := f(x)$, learn $y := x + f(x)$
 - Allows the network to learn residual functions
 - Allows training of very deep networks
 - ResNet-152 has 152 layers (vs. 19 in VGG-19)
 - Learn correction to previous features
 - Each block learns to correct the input rather than transform it completely
 - Family of models: ResNet-18, ResNet-34, ResNet-50, ResNet-101
 - Different depths for different computational budgets
- <https://arxiv.org/abs/1512.03385>
- <https://arxiv.org/abs/1603.05027>

The Deep Learning Revolution

- ResNet (2015): Skip connections, ultra-deep networks



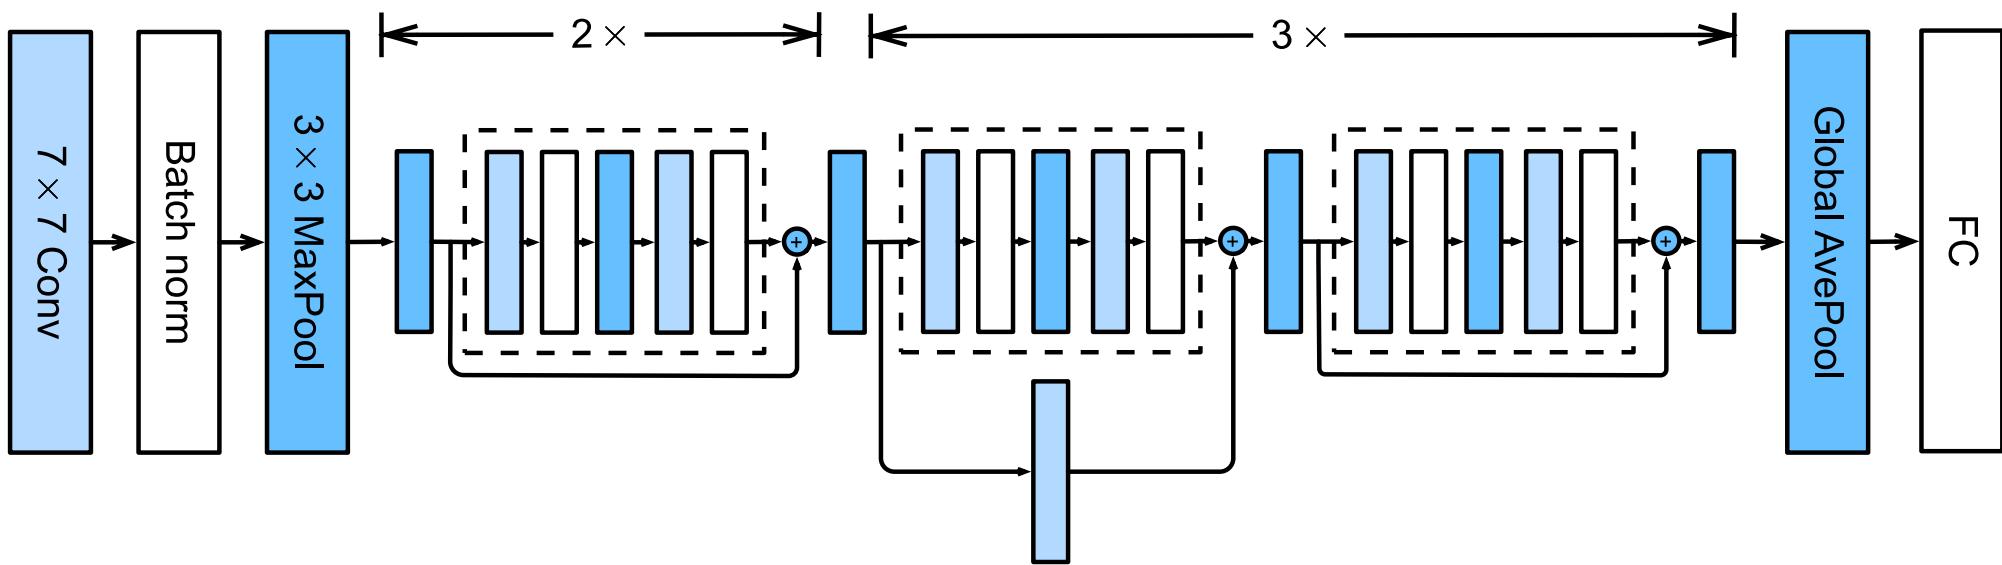
https://d2l.ai/chapter_convolutional-modern/resnet.html



https://d2l.ai/chapter_convolutional-modern/resnet.html

The Deep Learning Revolution

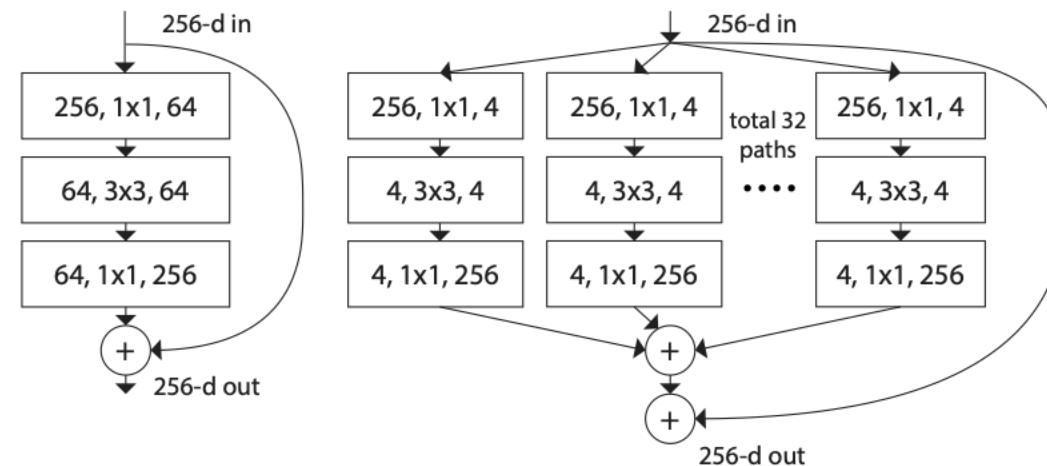
- ResNet (2015): ResNet-18



https://d2l.ai/chapter_convolutional-modern/resnet.html

The Deep Learning Revolution

- Variants of ResNet
 - ResNext: Information flows through several groups then aggregated
 - Combines ResNet with grouped convolutions (like inception block in GoogLeNet)
 - <https://arxiv.org/abs/1611.05431>



<https://arxiv.org/abs/1611.05431>

The Deep Learning Revolution

- Variants of ResNet
 - WideResNet: Wider blocks shown to be superior
 - Increases width instead of depth
 - <https://arxiv.org/abs/1605.07146>

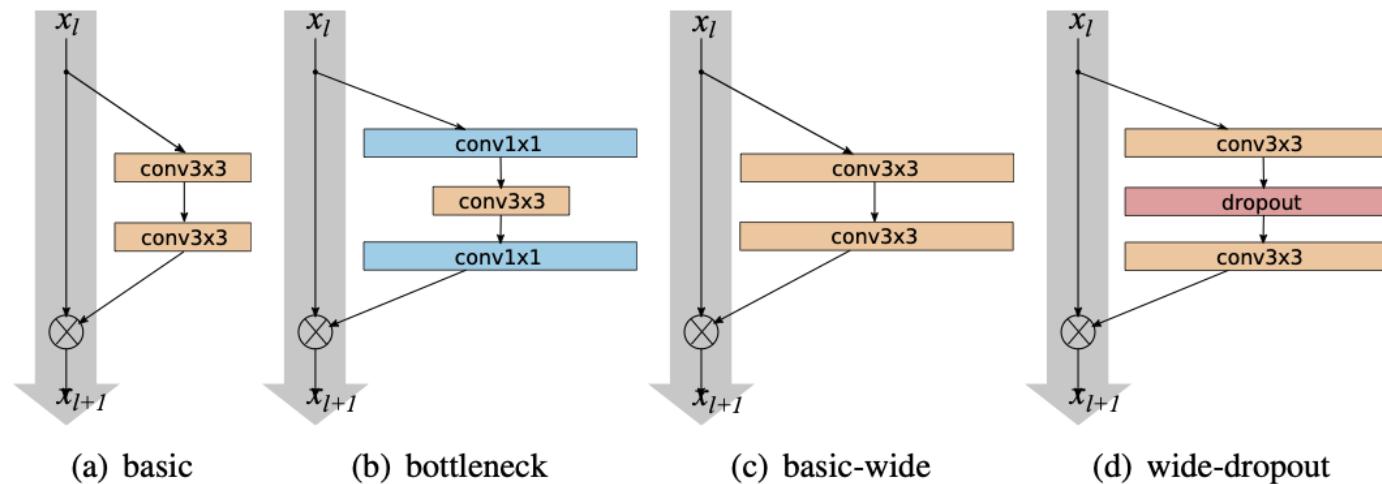


Figure 1: Various residual blocks used in the paper. Batch normalization and ReLU precede each convolution (omitted for clarity)

The Deep Learning Revolution

- DenseNet (2017): Dense connections between layers
 - Each layer is connected to all preceding layers
 - Improves gradient flow and feature reuse
 - <https://arxiv.org/abs/1608.06993>

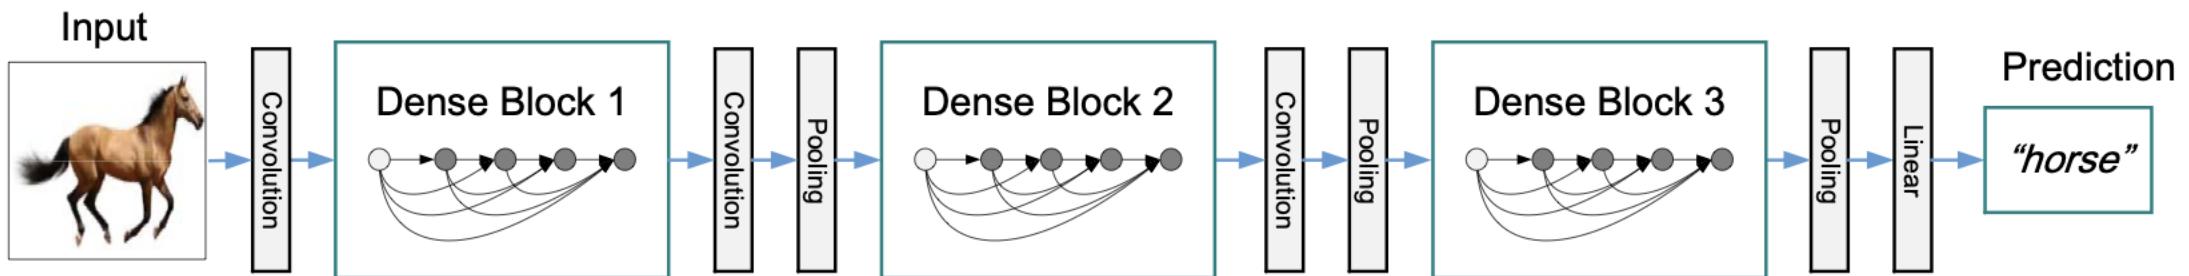


Figure 2: A deep DenseNet with three dense blocks. The layers between two adjacent blocks are referred to as transition layers and change feature-map sizes via convolution and pooling.

<https://arxiv.org/abs/1608.06993>

The Deep Learning Revolution

- DenseNet (2017): Dense connections between layers

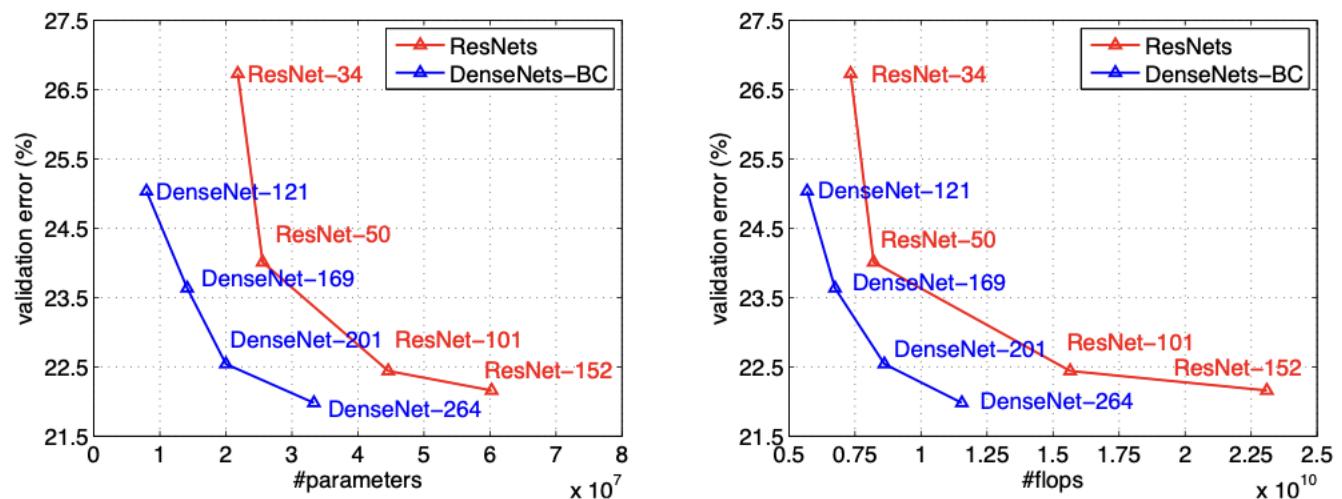


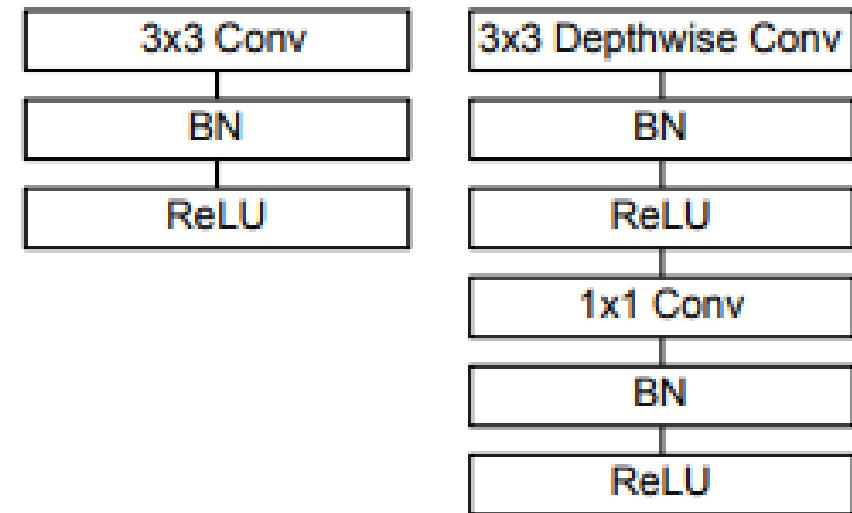
Figure 3: Comparison of the DenseNets and ResNets top-1 error rates (single-crop testing) on the ImageNet validation dataset as a function of learned parameters (*left*) and FLOPs during test-time (*right*).

The Deep Learning Revolution

- MobileNet (2017): Efficient CNNs for mobile devices
 - Uses 1×1 convolutions plus depthwise separable convolutions
 - Based on Canonical decomposition of a convolution kernel:
 - $V(x, y, t) = \sum_{i=x-\delta}^{x+\delta} \sum_{j=y-\delta}^{y+\delta} \sum_{s=1}^S K(i - x + \delta, j - y + \delta, s, t) U(i, j, s)$
 - $K(i, j, s, t) = \sum_{r=1}^R K_x(i, r) K_y(j, r) K_s(s, r) K_t(t, r)$
 - Converts many-dimensional kernel into several one-dimensional kernels
 - Depthwise convolutions treat each input channel independently
 - 1×1 convolution treat all channels of a single pixel
 - No pooling blocks
 - convolution with stride=2 used for downsampling
 - <https://arxiv.org/abs/1704.04861>

The Deep Learning Revolution

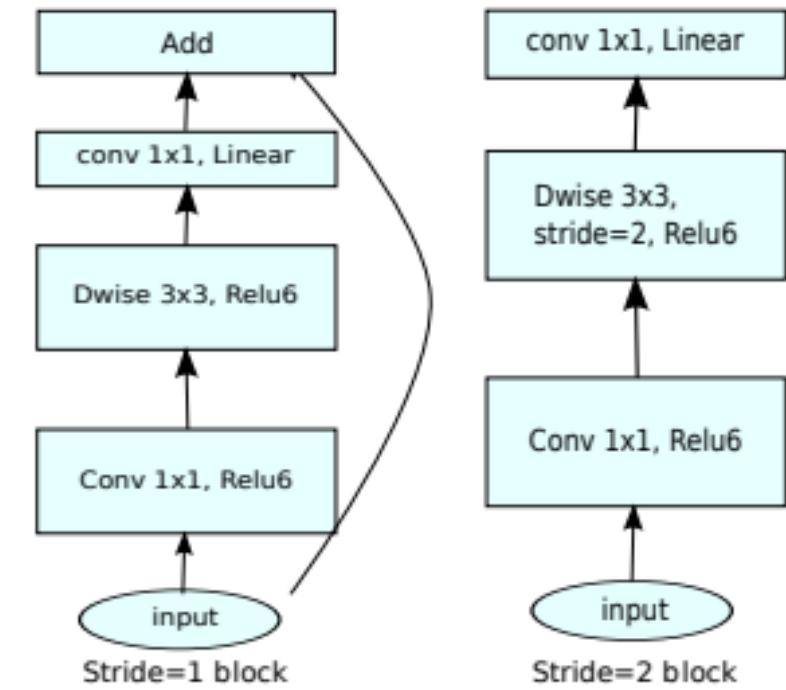
- MobileNet (2017): Efficient CNNs for mobile devices
 - Uses 1×1 convolutions plus depthwise separable convolutions
 - Based on Canonical decomposition of a convolution kernel:
 - No pooling blocks
 - <https://arxiv.org/abs/1704.04861>



<https://raw.githubusercontent.com/oseledets/dl2024/8d739d06892c9cc71f198c1c478d51b4a6801866/lectures/lecture-2/mobilenet-block.png>

The Deep Learning Revolution

- MobileNetV2 (2018): Inverted bottleneck design
 - First 1×1 convolution increases channels, last decreases
 - Expands channels in the middle of the block
 - Some blocks have residual connections, others do not
 - Combines ideas from ResNet and MobileNet
 - <https://arxiv.org/abs/1801.04381>



<https://raw.githubusercontent.com/oseledets/dl2024/8d739d06892c9cc71f198c1c478d51b4a6801866/lectures/lecture-2/mobilenetv2.png>

The Deep Learning Revolution

Network Architecture	Number of Parameters	Top-1 Accuracy	Top-5 Accuracy
Xception	22.91M	0.790	0.945
VGG16	138.35M	0.715	0.901
MobileNetV1 (alpha=1, rho=1)	4.20M	0.709	0.899
MobileNetV1 (alpha=0.75, rho=0.85)	2.59M	0.672	0.873
MobileNetV1 (alpha=0.25, rho=0.57)	0.47M	0.415	0.663
MobileNetV2 (alpha=1.4, rho=1)	6.06M	0.750	0.925
MobileNetV2 (alpha=1, rho=1)	3.47M	0.718	0.910
MobileNetV2 (alpha=0.35, rho=0.43)	1.66M	0.455	0.704

<https://github.com/oseledets/dl2024/blob/main/lectures/lecture-2/lecture-2.ipynb>

The Deep Learning Revolution

- EfficientNet (2019): Balanced scaling of network dimensions
 - All CNNs are similar and contain several stages
 - Scaling only a single parameter quickly saturates after 80% accuracy

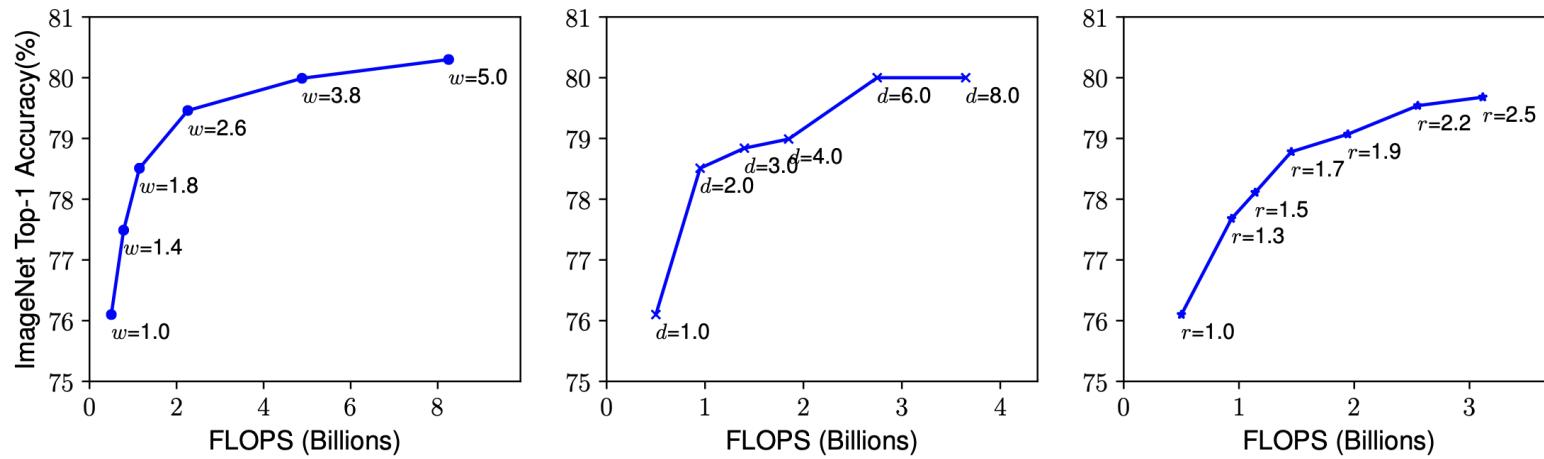


Figure 3. Scaling Up a Baseline Model with Different Network Width (w), Depth (d), and Resolution (r) Coefficients. Bigger networks with larger width, depth, or resolution tend to achieve higher accuracy, but the accuracy gain quickly saturates after reaching 80%, demonstrating the limitation of single dimension scaling. Baseline network is described in Table 1.

<https://arxiv.org/abs/1905.11946>

The Deep Learning Revolution

- EfficientNet (2019): Balanced scaling of network dimensions
 - Compound scaling
 - Scales all three dimensions together
 - Depth $d = \alpha^\phi, \alpha \geq 1$
 - Width $w = \beta^\phi, \beta \geq 1$
 - Resolution $r = \gamma^\phi, \gamma \geq 1$
 - Such that $\alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$
 - Defined by a grid search
 - <https://arxiv.org/abs/1905.11946>

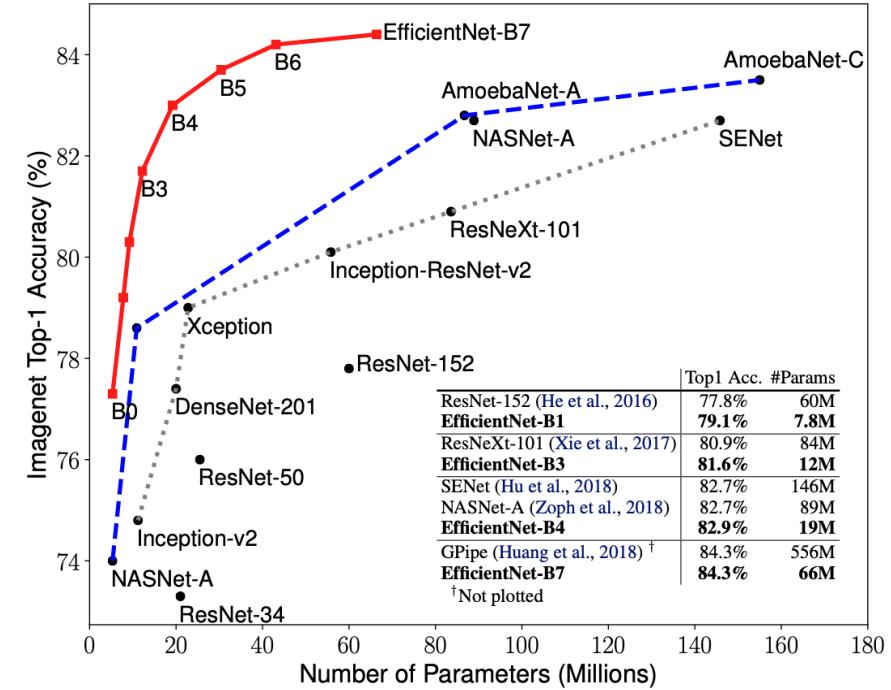


Figure 1. Model Size vs. ImageNet Accuracy. All numbers are for single-crop, single-model. Our EfficientNets significantly outperform other ConvNets. In particular, EfficientNet-B7 achieves new state-of-the-art 84.3% top-1 accuracy but being 8.4x smaller and 6.1x faster than GPipe. EfficientNet-B1 is 7.6x smaller and 5.7x faster than ResNet-152. Details are in Table 2 and 4.

<https://arxiv.org/abs/1905.11946>

Transformers

- Vision Transformers (2020): Incorporating transformer architecture
 - Treats images as sequences of patches
 - Uses self-attention instead of convolutions
 - Outperforms CNNs for many tasks
 - <https://arxiv.org/abs/2010.11929>